


Meta Learning: Learn to learn

Hung-yi Lee

What does “meta” mean? $\text{meta-X} = \text{X about X}$

Source of image: <https://medium.com/intuitionmachine/the-brute-force-method-of-deep-learning-innovation-58b497323ae5> (Denny Britz's graphic)


這門課的作業在做甚麼？




朋友覺得我在



我媽覺得我在




大眾覺得我在



指導教授覺得我在



我以為我在



事實上我在

感謝 沈昇勳 同學提供圖檔

Industry



Using 1000 GPUs to try
1000 sets of hyperparameters

Academia



“Telepathize” (通靈) a set of
good hyperparameters

Can machine automatically determine the hyperparameters?

Machine Learning 101

Machine Learning

= Looking for a function


Step 1: Function
with unknown

Step 2: Define
loss function

Step 3:
Optimization

Dog-Cat Classification

$$f(\text{image of cat}) = \text{"cat"}$$



Weights and biases of neurons are unknown parameters (*learnable*).


Using θ to represent the unknown parameters.

Machine Learning

Step 1: Function
with unknown

Step 2: Define
loss function





Step 3:
Optimization



$$L(\theta)$$



$$L(\theta) = \sum_{k=1}^K e_k$$

Training Examples



Cross-entropy e_1


e_2




Ground Truth

Machine Learning 101

Step 1: Function
with unknown



Step 2: Define
loss function




Step 3:
Optimization

loss: $L(\boldsymbol{\theta}) = \sum_{k=1}^K e_k$ sum over
training
examples

$$\boldsymbol{\theta}^* = \underset{\boldsymbol{\theta}}{\operatorname{arg\,min}} L(\boldsymbol{\theta})$$

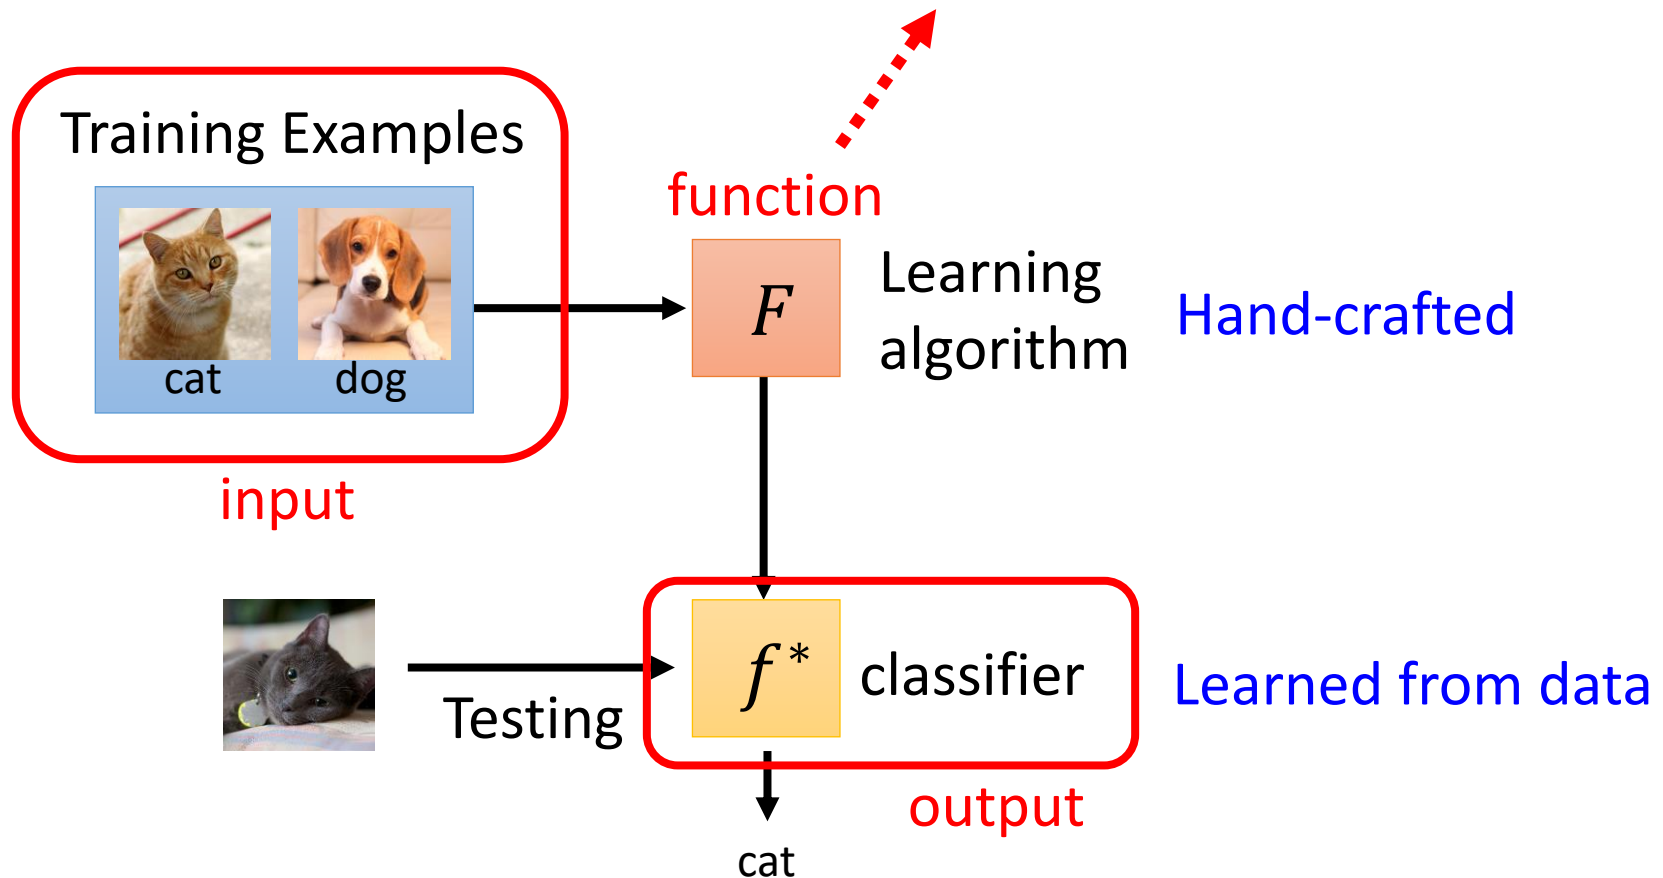
done by gradient descent

$f_{\boldsymbol{\theta}^*}$ is the function learned by
learning algorithm from data



Introduction of Meta Learning

What is Meta Learning?



Meta Learning – Step 1

- What is *learnable* in a learning algorithm?

Training Examples



F

Deep Learning

Component

Net Architecture,
Initial Parameters,
Learning Rate,
.....



Testing

f^*

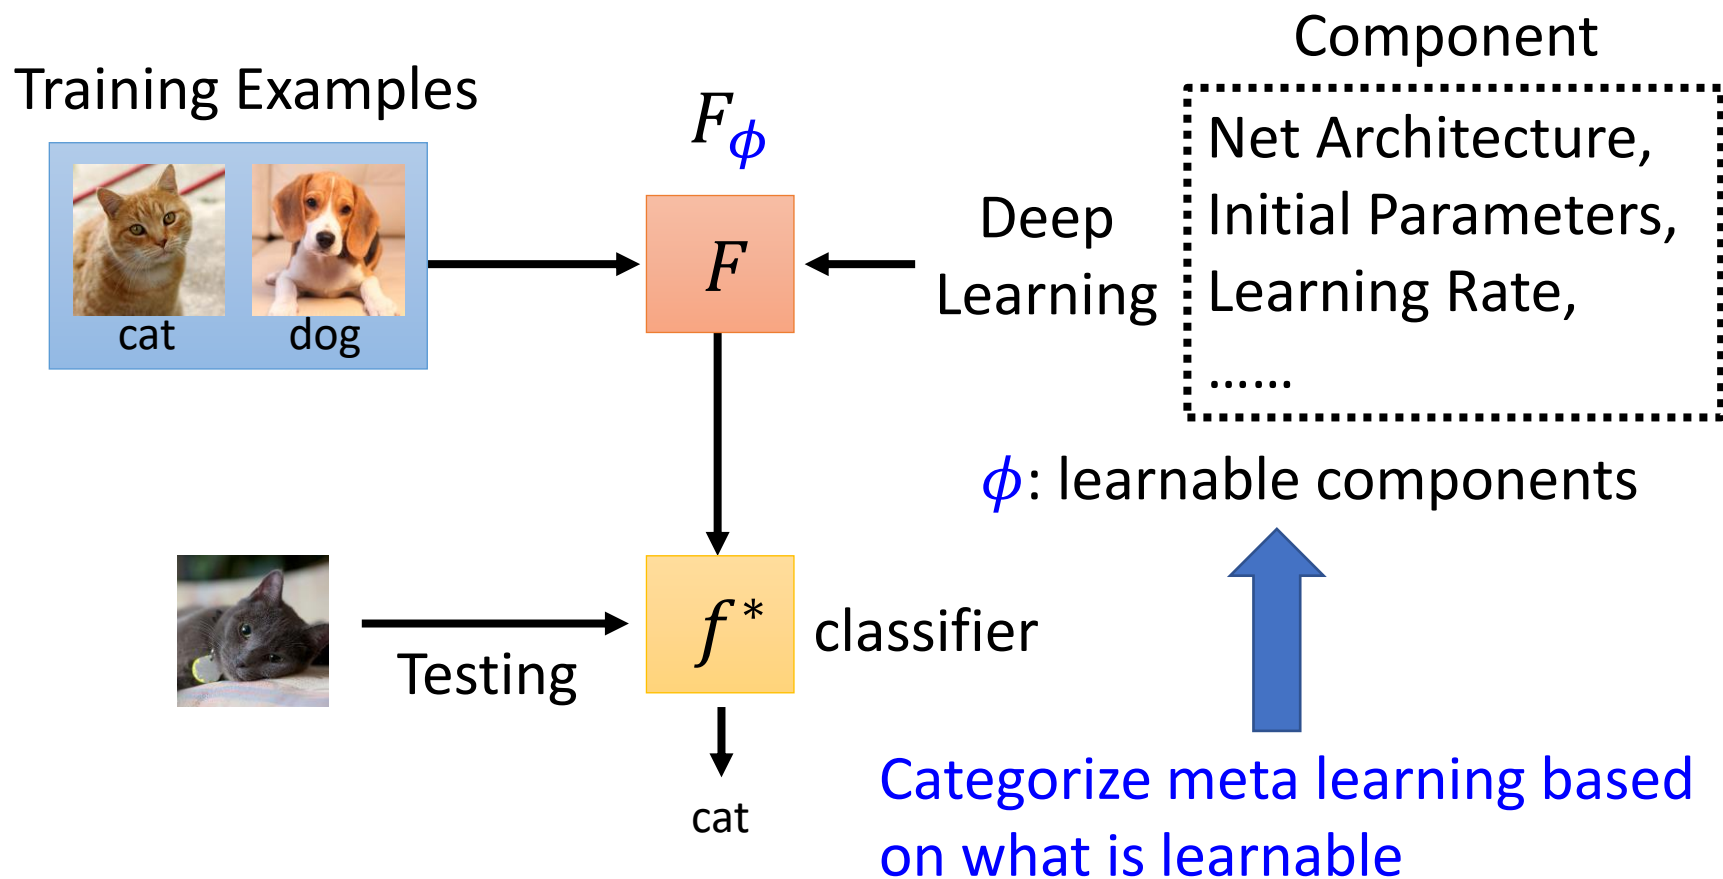
classifier

cat

In meta, we will try to
learn some of them.

Meta Learning – Step 1

- What is **learnable** in a learning algorithm?



Meta Learning – Step 2


- Define loss function for learning algorithm F_ϕ
 $L(\phi)$




Training Tasks

Task 1
Apple &
Orange

Train




Test




Task 2
Car & Bike

Train






Test



Meta Learning – Step 2

Task 1

*Training
Examples*



How to define $L(\phi)$




$L(\phi)$ ↓

θ^{1*} : parameters of the classifier learned by F_ϕ
using the training examples of task 1

Meta Learning – Step 2

Task 1

*Training
Examples*



classifier

How can we know a classifier is good or bad?

Evaluate the classifier on testing set


How to define $L(\phi)$

$L(\phi)$ ↑


Meta Learning – Step 2

Task 1

Training Examples



Testing Examples



F_ϕ

$f_{\theta^{1*}}$

prediction


l^1 Compute difference

Testing Examples



$f_{\theta^{1*}}$

$f_{\theta^{1*}}$



apple orange

apple orange

Cross-entropy

Cross-entropy



apple orange


apple orange

Ground Truth


Meta Learning – Step 2

Task 1

Training Examples



Testing Examples



F_ϕ



$f_{\theta^{1*}}$



prediction


l^1

Compute difference

Testing Examples



$f_{\theta^{1*}}$



apple orange



$f_{\theta^{1*}}$



apple orange

Cross-entropy

Cross-entropy



apple orange




apple orange

Ground Truth


Meta Learning – Step 2

Task 1

Training Examples



Testing Examples



F_ϕ



$f_{\theta^{1*}}$



prediction

Compute difference




l^1

Testing Examples



$f_{\theta^{1*}}$

$f_{\theta^{1*}}$



apple orange

apple orange

Cross-entropy

Cross-entropy



apple orange


apple orange

Ground Truth

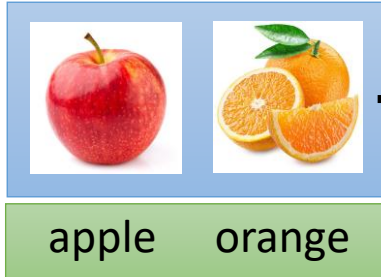
Meta Learning – Step 2

Task 1

*Training
Examples*



*Testing
Examples*




F_ϕ

$f_{\theta^{1*}}$


prediction

l^1

Task 2



*Testing
Examples*



F_ϕ

$f_{\theta^{2*}}$

prediction


l^2

Total loss: $L(\phi) = l^1 + l^2$ (sum over all the training tasks)


Meta Learning – Step 2

Task 1

Training Examples



Testing Examples




F_ϕ

$f_{\theta^{1*}}$


prediction

l^1

Task 2



Testing Examples



F_ϕ

$f_{\theta^{2*}}$

prediction

l^2

Total loss: $L(\phi) = \sum_{n=1}^N l^n$ (N is the number of the training tasks)


Meta Learning – Step 2

Task 1



In typical ML, you compute the loss based on **training examples**

In meta, you compute the loss based on **testing examples**

Hold on! You use **testing examples** during training???



Testing Examples



apple orange

apple orange

Ground Truth



Meta Learning – Step 2

Task 1

In typical ML, you compute the loss based on **training examples**

In meta, you compute the loss based on **testing examples** of **training tasks**.

Testing Examples



Meta Learning – Step 3

- Loss function for learning algorithm $L(\phi) = \sum_{n=1}^N l^n$
- Find ϕ that can minimize $L(\phi)$ $\phi^* = \arg \min_{\phi} L(\phi)$

- Using the optimization approach you know

If you know how to compute $\partial L(\phi) / \partial \phi$

Gradient descent is your friend.

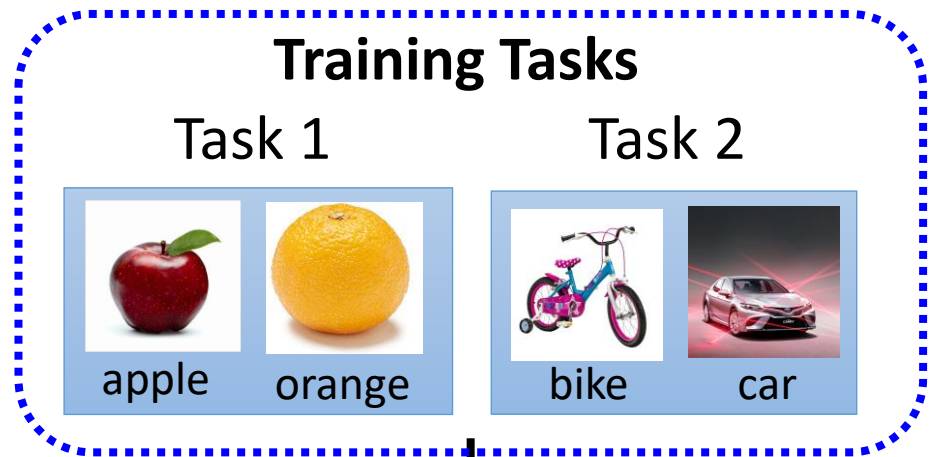
What if $L(\phi)$ is not differentiable?

Reinforcement Learning / Evolutionary Algorithm

Now we have a learned “learning algorithm” F_{ϕ^*}

Framework

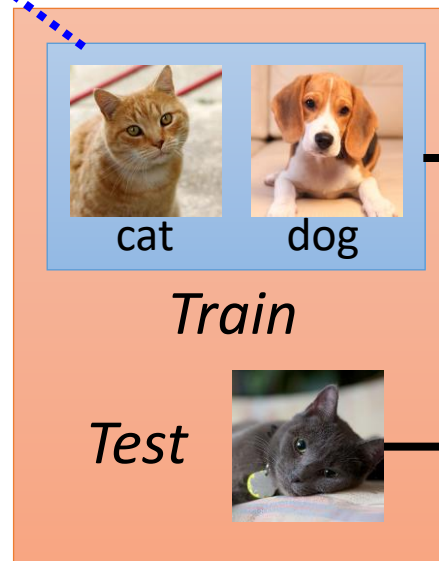
Not related to
the testing task



only need little labeled training data

**Testing
Task**

What we really
care about



F_{ϕ^*}

Learned
“Learning
Algorithm”

f_{θ^*}

cat

ML v.s. Meta

Goal

Machine Learning \approx find a function f

Dog-Cat
Classification

$$f(\text{img}) = \text{"cat"}$$

Meta Learning

\approx find a function F that finds a function f

Learning
Algorithm

$F($




cat



dog



cat



dog

Training Examples

$) = f$

Training Data

Machine Learning

One task




Train

Meta Learning


Training tasks

Task 1
Apple &
Orange

Train




Test




Task 2
Car & Bike

Train



Test




Support set

Query set

(in the literature of “learning to compare”)

Machine Learning

Within-task Training




Meta Learning

Training Tasks



Task 1

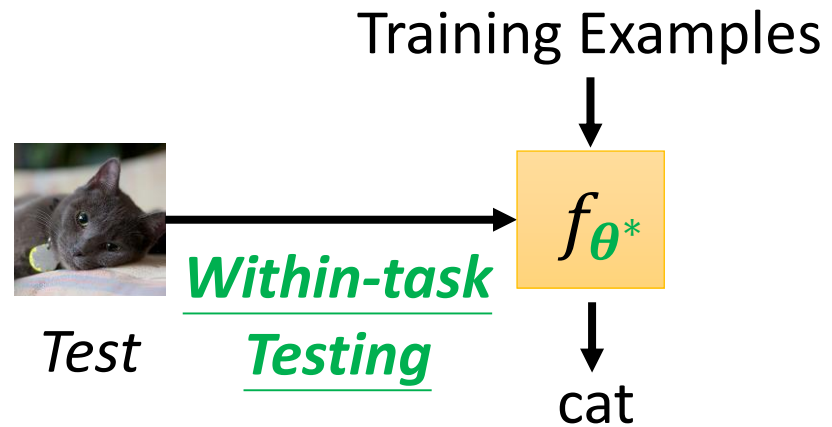
Task 2




Across-task Training

F_{ϕ^*} Learning Algorithm

Machine Learning



Meta Learning

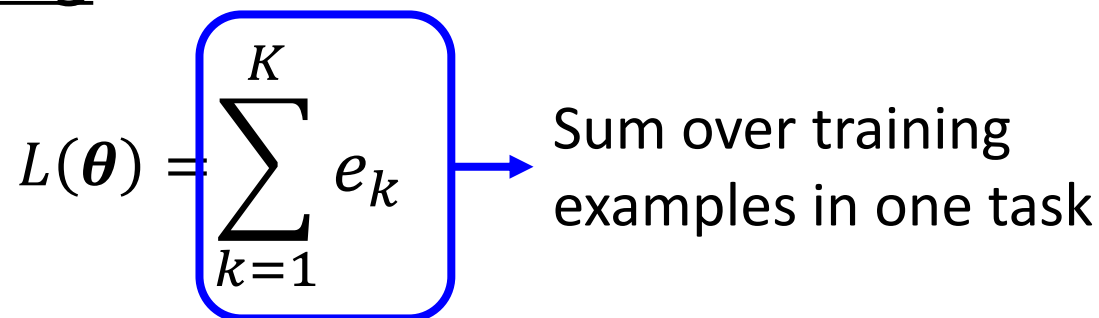


Loss

Machine Learning

$$L(\boldsymbol{\theta}) = \sum_{k=1}^K e_k$$

Sum over training examples in one task

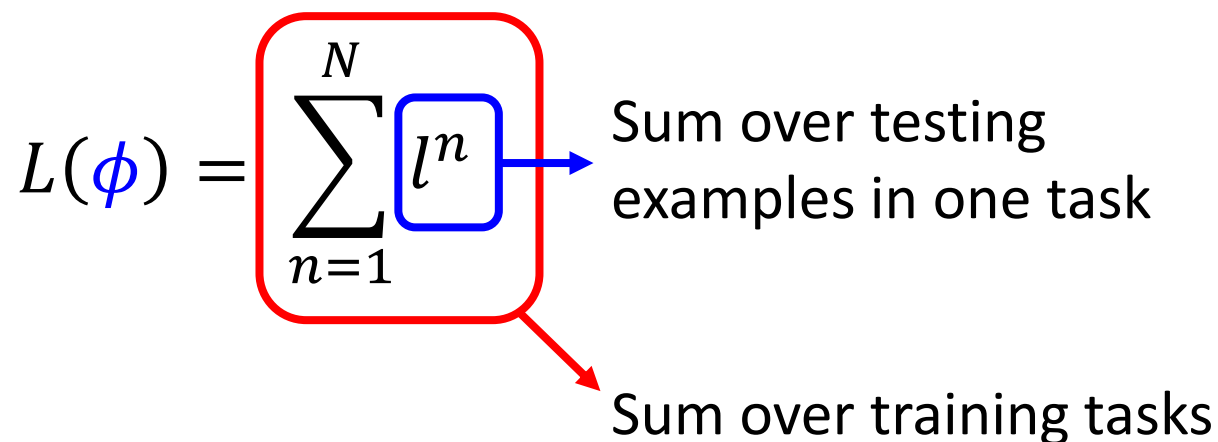


Meta Learning

$$L(\phi) = \sum_{n=1}^N l^n$$

Sum over testing examples in one task

Sum over training tasks



$$L(\phi) = \sum_{n=1}^N l^n$$

If your optimization method needs to compute $L(\phi)$

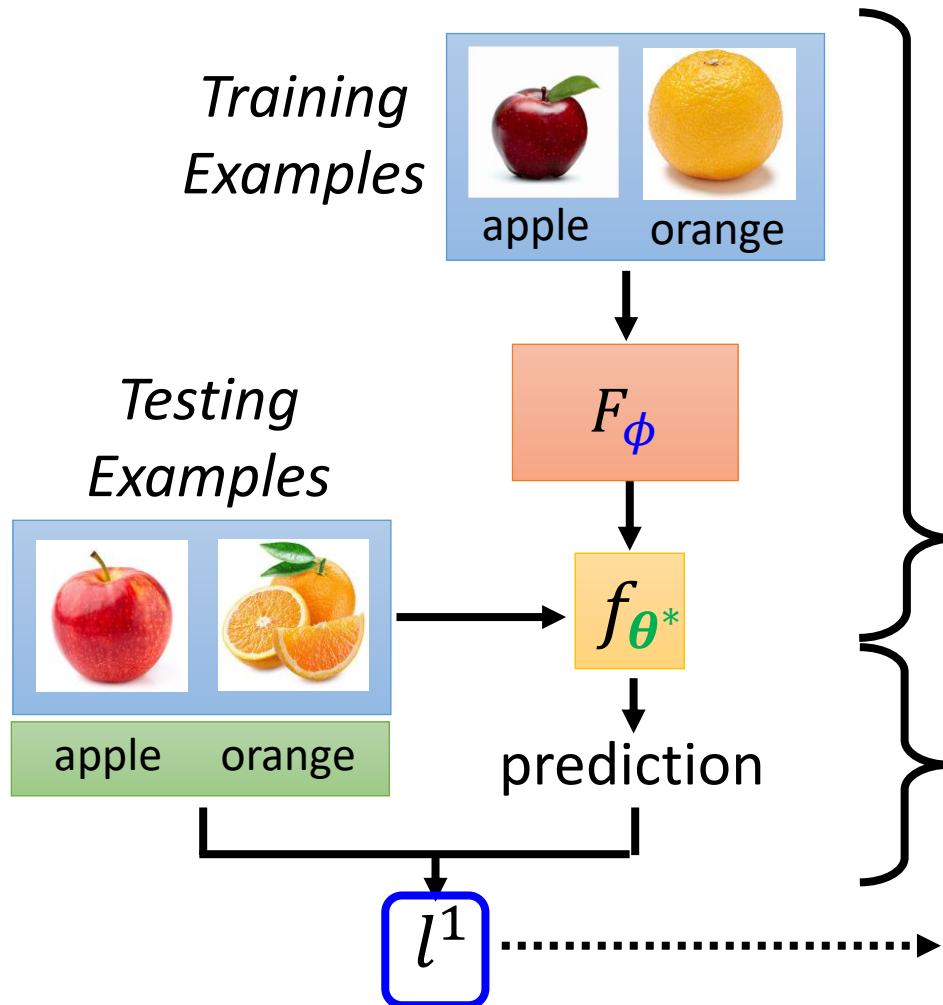
Outer Loop in
"Learning to initialize"

Across-task training
includes within-task
training and testing

Inner Loop in
"Learning to initialize"

Within-task Training

Within-task Testing



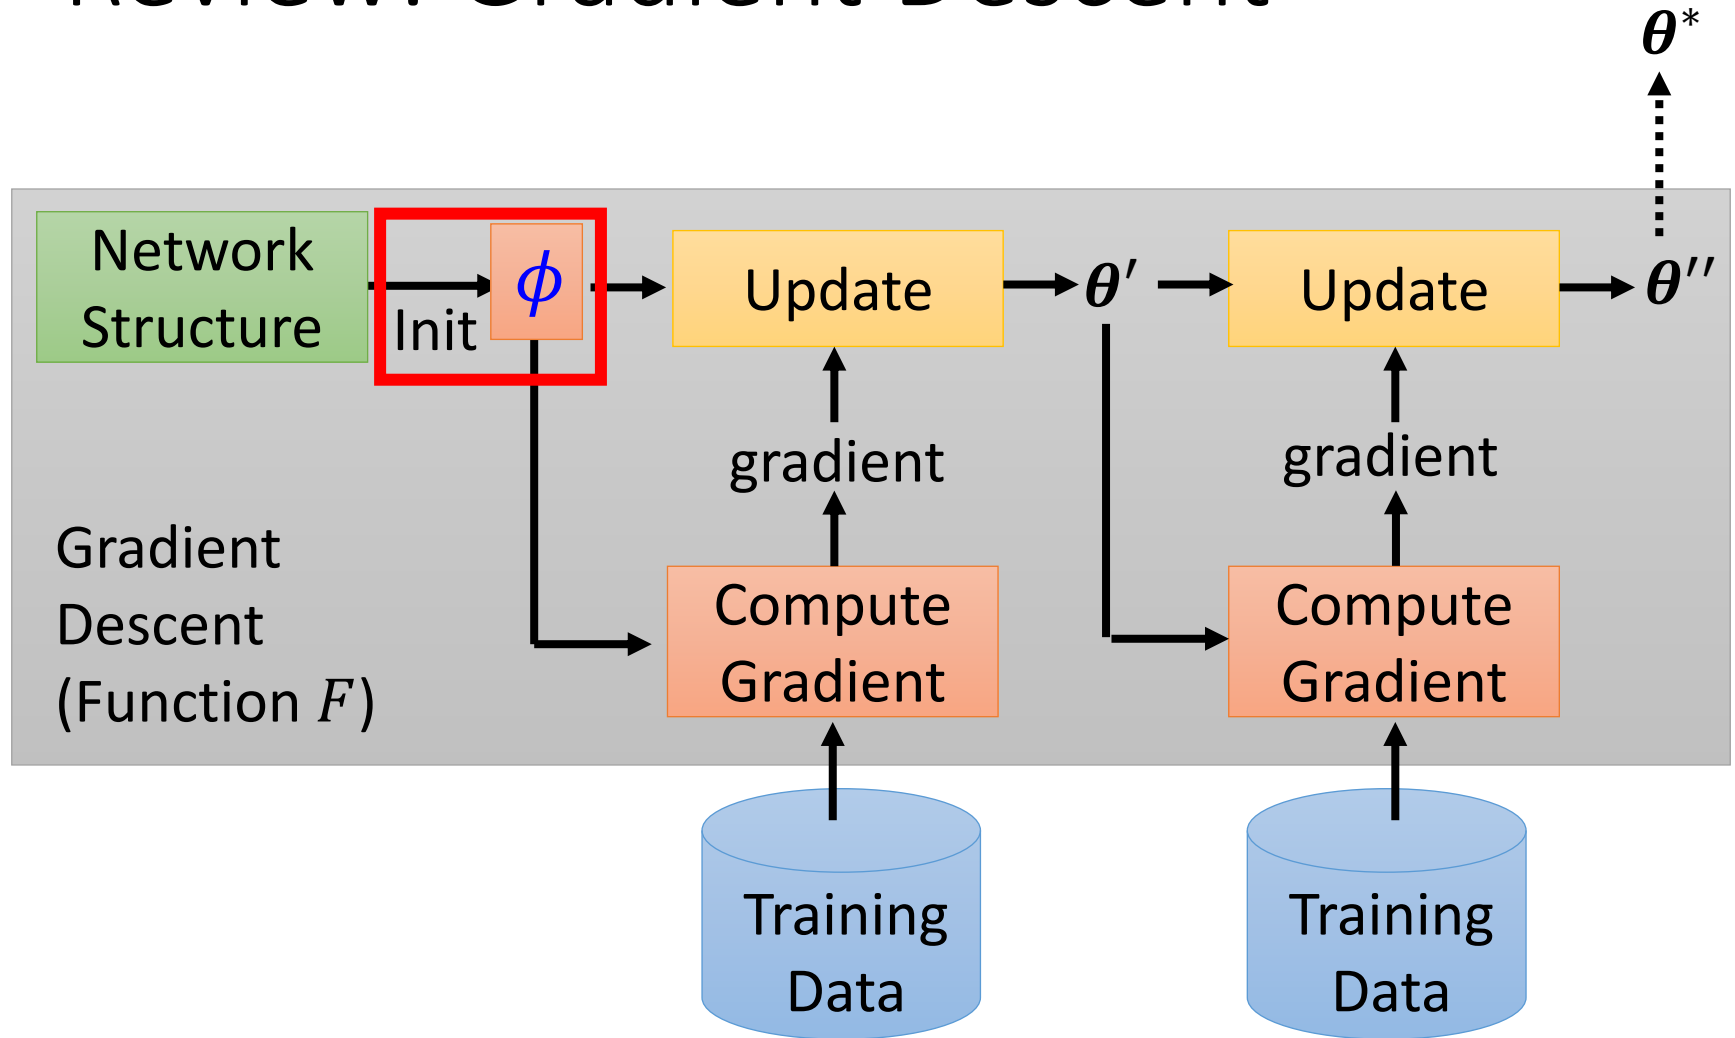
To compute the loss

Meta Learning v.s ML

- What you know about ML can usually apply to meta learning
 - Overfitting on training tasks
 - Get more training tasks to improve performance
 - Task augmentation
 - There are also hyperparameters when learning a learning algorithm
 - Development task 😊

What is learnable in a learning algorithm?

Review: Gradient Descent



Learning to initialize

- Model-Agnostic Meta-Learning (MAML)



Mammals


Chelsea Finn, Pieter Abbeel, and Sergey Levine, “Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks”, ICML, 2017

- Reptile



<https://arxiv.org/abs/1803.02999>

How to train your ~~Dragon~~ MAML



MAML

Task 1



cat



dog

Task 2



cat



dog



find good init

Testing
Task



cat



dog

Pre-training (Self-supervised Learning)



find good init



cat



dog

Trained by proxy tasks
(fill-in the blanks, etc.)

MAML

Isn't it domain adaptation / transfer learning?

Task 1



cat



dog

Task 2



cat



dog



find good init



cat



dog

Pre-training (more typical ways)



cat



dog



cat



dog



find good init



cat



dog

Use data from different tasks
to train a model

Also known as multi-task
learning (baseline of meta)

MAML v.s. Pre-training

- https://youtu.be/vUwOA3SNb_E



影片中有防不勝防
的業配

這就是 “meta 業配”



MAML is good because

- ANIL (Almost No Inner Loop)



Aniruddh Raghu, Maithra Raghu, Samy Bengio, Oriol Vinyals, Rapid Learning or Feature Reuse? Towards Understanding the Effectiveness of MAML, ICLR, 2020

More about MAML

- More mathematical details behind MAML
 - https://youtu.be/mxqzGwP_Qys
- First order MAML (FOMAML)
 - <https://youtu.be/3z997JhL9Oo>
- Reptile
 - <https://youtu.be/9jJe2AD35P8>

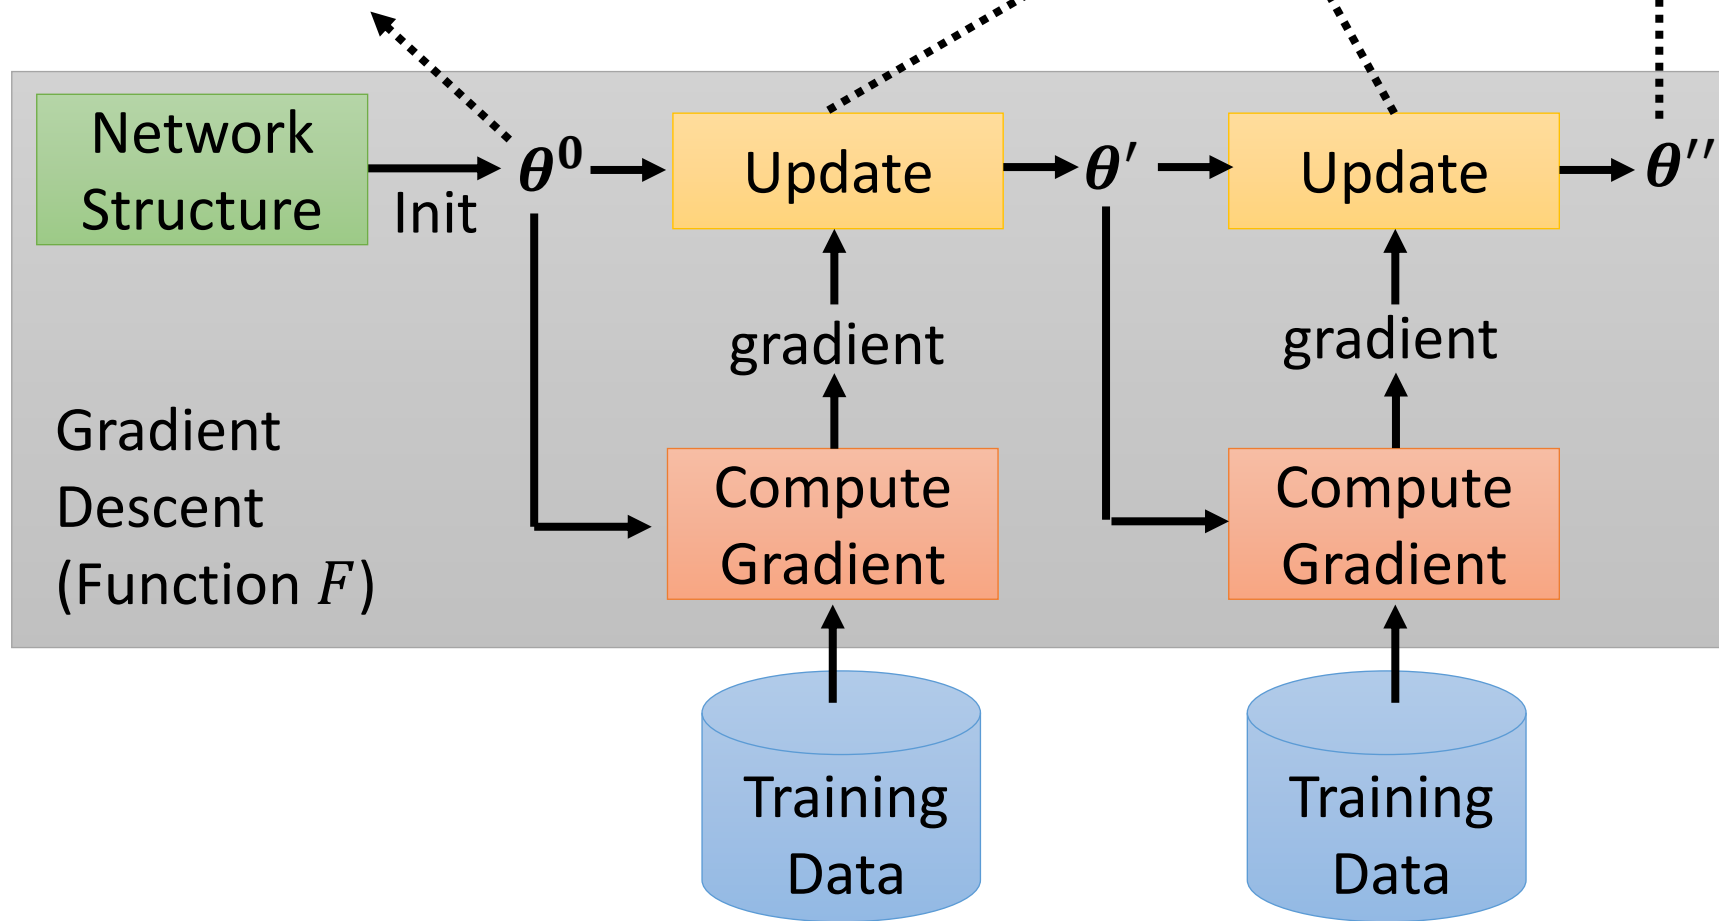
Optimizer

Basis form: $\theta^{t+1} \leftarrow \theta^t - \lambda g^t$

Adagrad, RMSprop, NAG, Adam

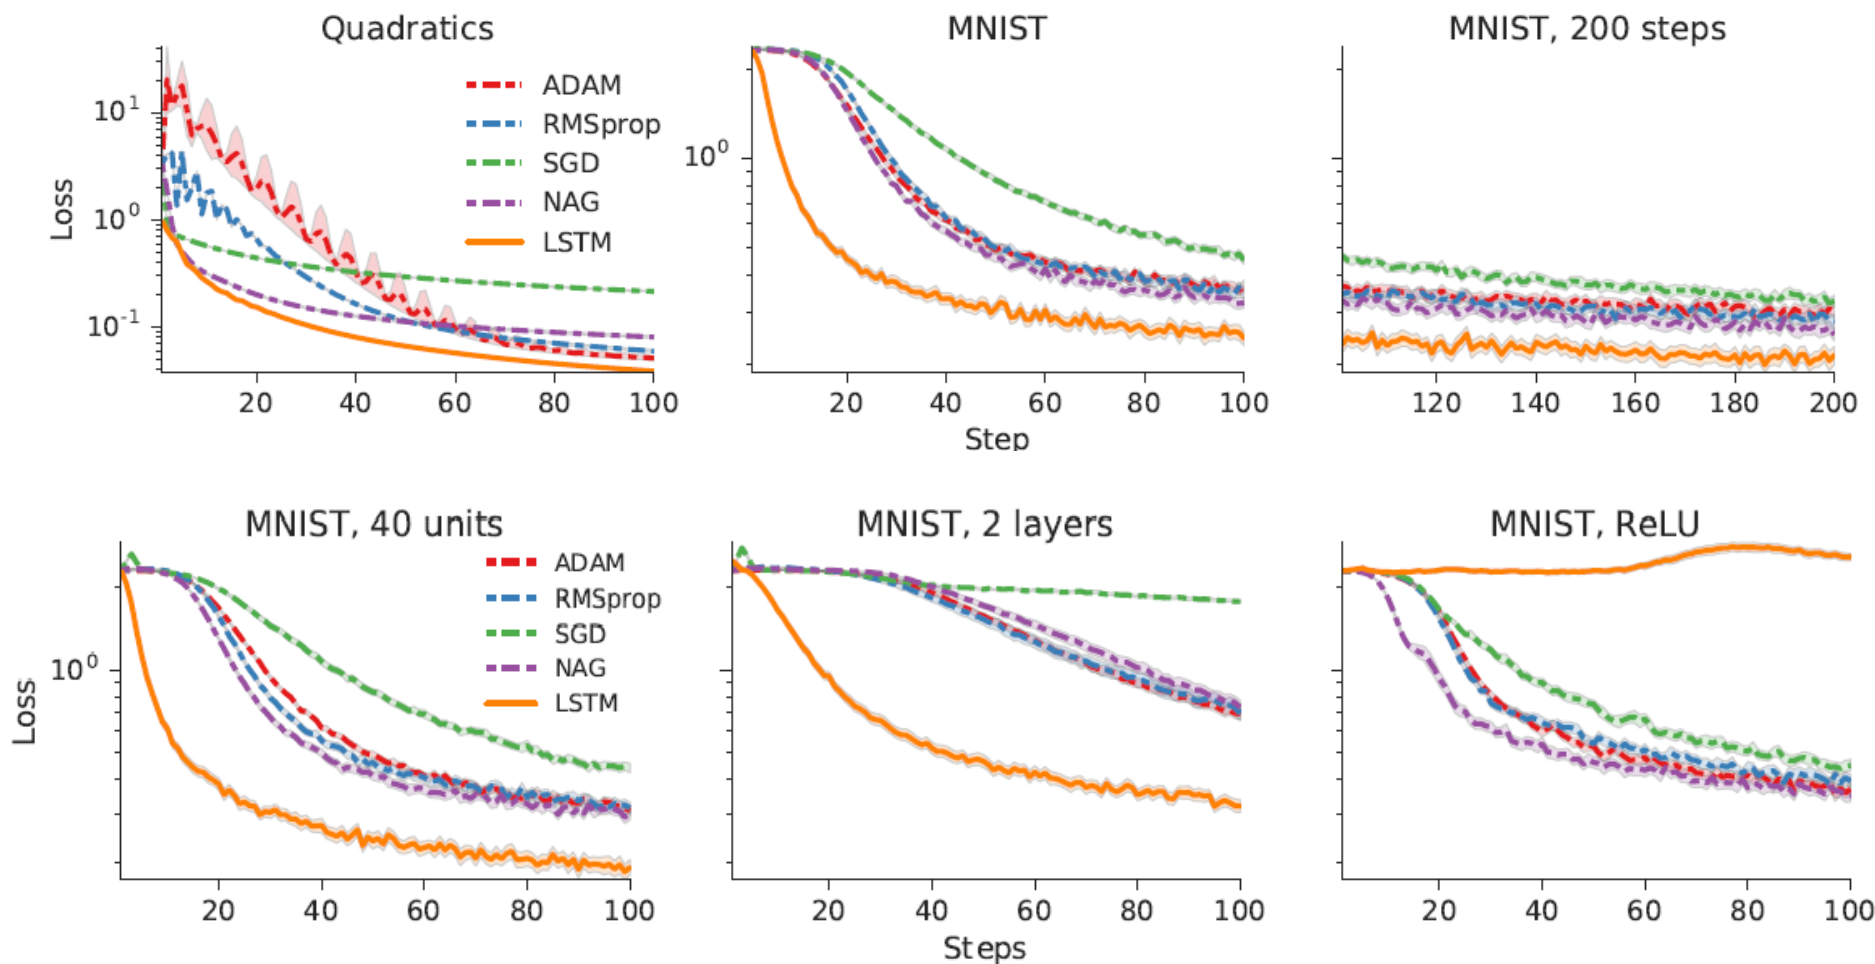
Is the optimizer learnable?

Can be learned by MAML

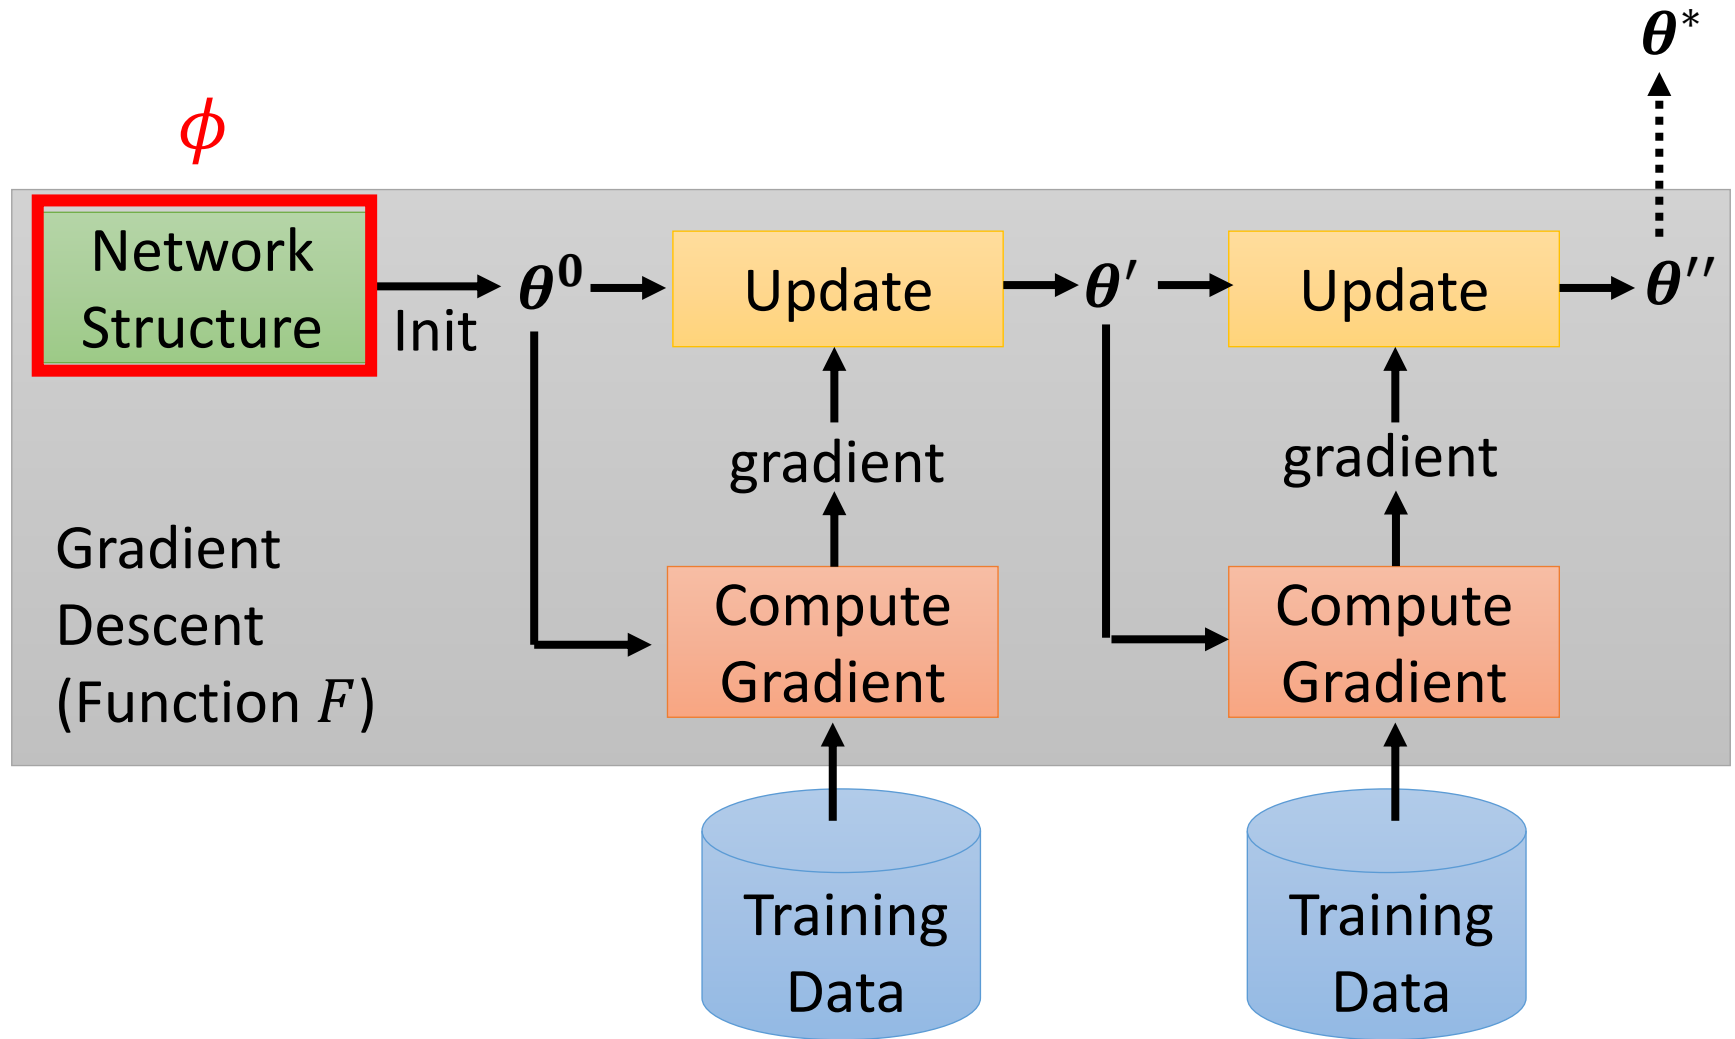


Optimizer

Marcin Andrychowicz, et al., Learning to learn by gradient descent by gradient descent, NIPS, 2016



Network Architecture Search (NAS)



Network Architecture Search (NAS)

$$\hat{\phi} = \arg \min_{\phi} L(\phi) \quad \nabla_{\phi} L(\phi) = ?$$

 Network
Architecture

- Reinforcement Learning

- Barret Zoph, et al., Neural Architecture Search with Reinforcement Learning, ICLR 2017
- Barret Zoph, et al., Learning Transferable Architectures for Scalable Image Recognition, CVPR, 2018
- Hieu Pham, et al., Efficient Neural Architecture Search via Parameter Sharing, ICML, 2018

An agent uses a set of actions to
determine the network architecture.

ϕ : the agent's parameters


$$-L(\phi)$$

Reward to be
maximized

Network Architecture Search (NAS)

Across-task
Training


Update ϕ to maximize reward $-L(\phi)$



$-L(\phi)$

Accuracy
of the
network


form a
network



Train the network

Within-task Training

Network Architecture Search (NAS)

$$\hat{\phi} = \arg \min_{\phi} L(\phi) \quad \nabla_{\phi} L(\phi) = ?$$


Network
Architecture

- Reinforcement Learning

- Barret Zoph, et al., Neural Architecture Search with Reinforcement Learning, ICLR 2017
- Barret Zoph, et al., Learning Transferable Architectures for Scalable Image Recognition, CVPR, 2018
- Hieu Pham, et al., Efficient Neural Architecture Search via Parameter Sharing, ICML, 2018

- Evolution Algorithm


- Esteban Real, et al., Large-Scale Evolution of Image Classifiers, ICML 2017
- Esteban Real, et al., Regularized Evolution for Image Classifier Architecture Search, AAAI, 2019
- Hanxiao Liu, et al., Hierarchical Representations for Efficient Architecture Search, ICLR, 2018

Network Architecture Search (NAS)


$$\hat{\phi} = \arg \min_{\phi} L(\phi) \quad \nabla_{\phi} L(\phi) = ?$$

Network
Architecture

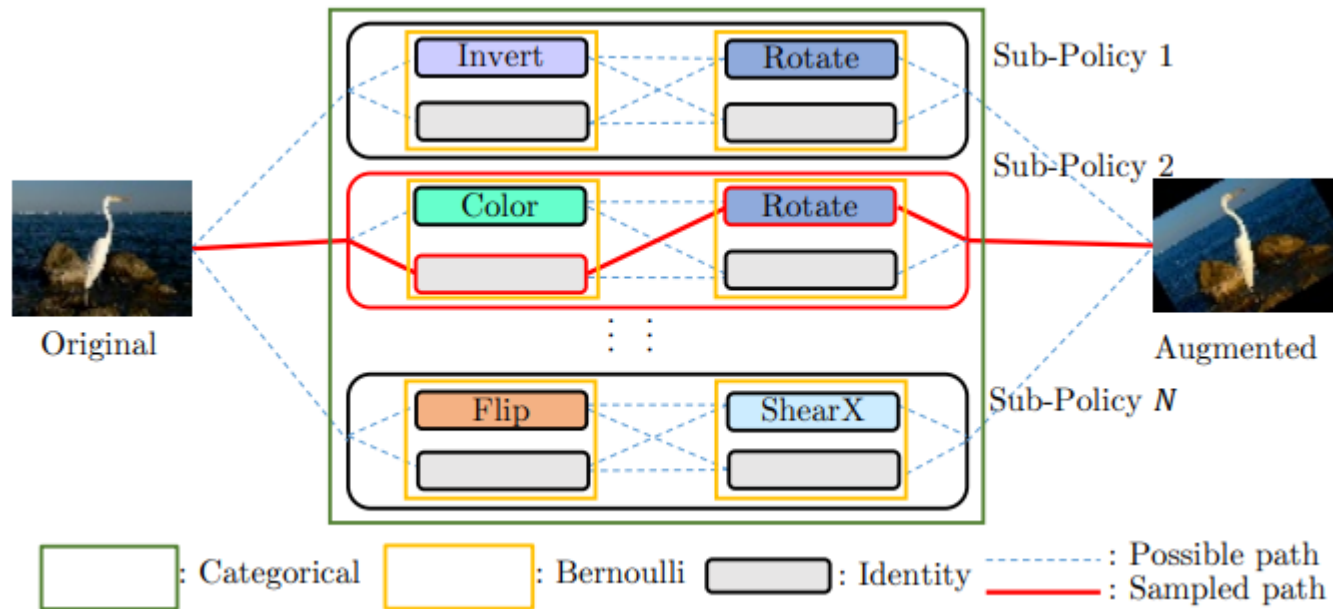
- DARTS Hanxiao Liu, et al., DARTS: Differentiable Architecture Search, ICLR, 2019



Data Processing?



Data Augmentation



Yonggang Li, Guosheng Hu, Yongtao Wang, Timothy Hospedales, Neil M. Robertson, Yongxin Yang, DADA: Differentiable Automatic Data Augmentation, ECCV, 2020

Daniel Ho, Eric Liang, Ion Stoica, Pieter Abbeel, Xi Chen, Population Based Augmentation: Efficient Learning of Augmentation Policy Schedules, ICML, 2019


Ekin D. Cubuk, Barret Zoph, Dandelion Mane, Vijay Vasudevan, Quoc V. Le, AutoAugment: Learning Augmentation Policies from Data, CVPR, 2019

Sample Reweighting

- Give different samples different weights

Larger weights (focus on tough examples)?

Smaller weights (the labels are noisy)?



Sample Weighting Strategies ➡ Learnable ϕ

Jun Shu, Qi Xie, Lixuan Yi, Qian Zhao, Sanping Zhou, Zongben Xu, Deyu Meng, Meta-Weight-Net: Learning an Explicit Mapping For Sample Weighting, NeurIPS, 2019
Mengye Ren, Wenyuan Zeng, Bin Yang, Raquel Urtasun, Learning to Reweight Examples for Robust Deep Learning, ICML, 2018



Beyond Gradient Descent

Andrei A. Rusu, Dushyant Rao, Jakub Sygnowski, Oriol Vinyals, Razvan Pascanu, Simon Osindero, Raia Hadsell,
Meta-Learning with Latent Embedding Optimization, ICLR, 2019

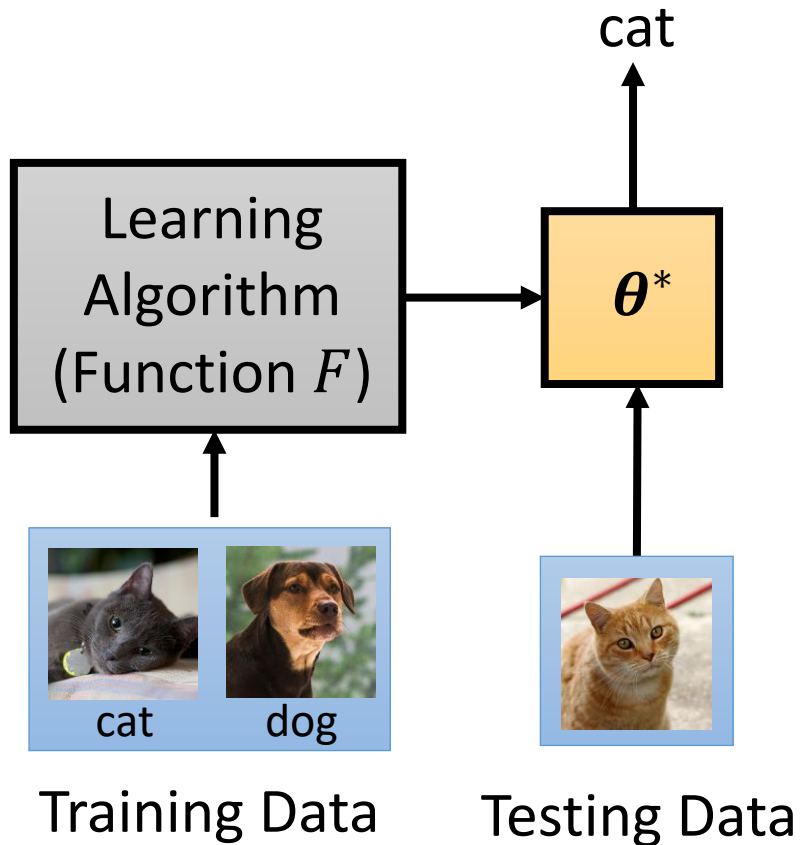
θ^*

This is a Network.
Its parameter is ϕ

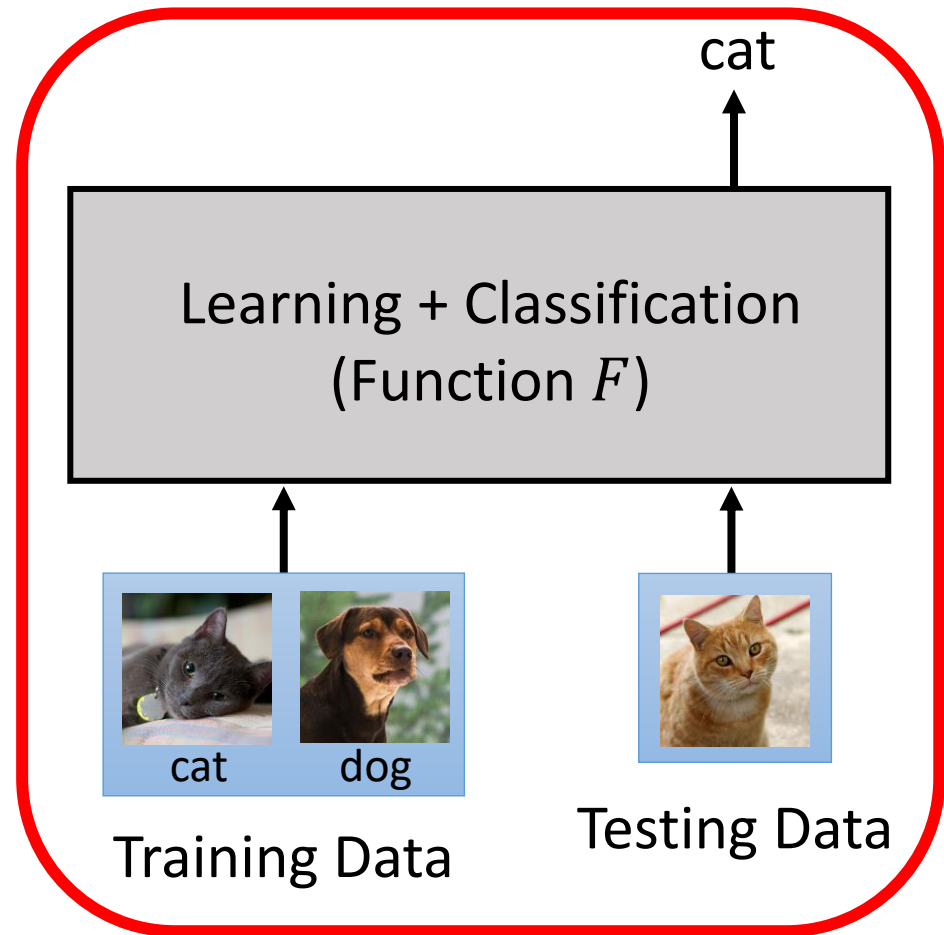
(Invent new learning algorithm! Not gradient descent anymore)



Until now



How about?



Learning to compare
(metric-based approach)

https://youtu.be/yyKaACh_j3M

<https://youtu.be/scK2EIT7klw>

<https://youtu.be/semSxPP2Yzg>

https://youtu.be/ePimv_k-H24



Applications

Few-shot Image Classification

- Each class only has a few images.



Class 1



Class 1



Class 2



Class 2



Class 3



Class 3



Which
class?


3-ways 2-shot

- **N-ways K-shot** classification: In each task, there are **N classes**, each has **K examples**.
- In meta learning, you need to prepare many **N-ways K-shot** tasks as training and testing tasks.

Omniglot

<https://github.com/brendenlake/omniglot>

- 1623 characters
- Each has 20 examples



Omniglot

Demo:

<https://openai.com/blog/reptile/>

20 ways

1 shot

Each character
represents a class

ᱠ	ᱡ	ᱢ	ᱣ	ᱤ
ᱥ	ᱦ	ᱧ	ᱨ	ᱩ
ᱪ	ᱫ	ᱬ	ᱭ	ᱮ
ᱯ	ᱰ	ᱱ	ᱲ	ᱳ



Testing set
(Query set)

Training set
(Support set)

- Split your characters into training and testing characters
 - Sample N training characters, sample K examples from each sampled characters → one training task
 - Sample N testing characters, sample K examples from each sampled characters → one testing task

	(A) Learning to initialize	(B) Learning to compare	(C) Other
Sound Event Detection	(Shi et al., 2020)	(Shimada et al., 2020a) (Chou et al., 2019) (Wang et al., 2020) (Shimada et al., 2020b) (Shi et al., 2020)	Network architecture search: (Li et al., 2020)
Keyword Spotting	(Chen et al., 2020a)	(Huh et al., 2020)	Net2Net: (Veniat et al., 2019) Network architecture search: (Mazzawi et al., 2019) Network architecture search: (Mo et al., 2020)
Text Classification	(Dou et al., 2019) (Bansal et al., 2019)	(Yu et al., 2018) (Tan et al., 2019) (Geng et al., 2019) (Sun et al., 2019)	Learning the learning algorithm: (Wu et al., 2019)
Voice Cloning			Learning the learning algorithm: (Chen et al., 2019b) (Serrà et al., 2019)
Sequence Labeling	(Wu et al., 2020)	(Hou et al., 2020)	
Machine Translation	(Gu et al., 2018) (Indurthi et al., 2020)		
Speech Recognition	(Hsu et al., 2020) (Klejch et al., 2019) (Winata et al., 2020a) (Winata et al., 2020b)		Learning to optimize: (Klejch et al., 2018) Network architecture search: (Chen et al., 2020b) (Baruwa et al., 2019)
Knowledge Graph	(Obamuyide and Vlachos, 2019) (Bose et al., 2019) (Lv et al., 2019) (Wang et al., 2019)	(Ye and Ling, 2019) (Chen et al., 2019a) (Xiong et al., 2018) (Gao et al., 2019)	
Dialogue / Chatbot	(Qian and Yu, 2019) (Madotto et al., 2019) (Mi et al., 2019)		Learning to optimize: (Chien and Licow, 2019)
Parsing	(Guo et al., 2019) (Huang et al., 2018)		
Word Embedding	(Hu et al., 2019)	(Sun et al., 2018)	
Multi-model		(Eloff et al., 2019)	Learning the learning algorithm: (Surís et al., 2019)

http://speech.ee.ntu.edu.tw/~tlkagk/meta_learning_table.pdf